

Bayesian Hierarchical Dynamic Factor Models

Anthony M. Thomas, Jr.
anthony.thomas3@mavs.uta.edu

Department of Mathematics
The University of Texas at Arlington

August 17, 2020

Contents

- ① Background
- ② Variational Bayesian Inference
- ③ Hierarchical Dynamic Factor Analysis
 - Classical Factor Analysis
 - Dynamic Factor Analysis

Section 1

- 1 Background
- 2 Variational Bayesian Inference
- 3 Hierarchical Dynamic Factor Analysis
 - Classical Factor Analysis
 - Dynamic Factor Analysis

Factor Analysis

- ① **Factor Analysis** is a method that uses the covariances between a set of observed variables to described them in terms of a smaller set of unobservable variables called factors.

Bayesian Inference

Bayesian Inference can be described by two parts:

- ① Build a model based on data \mathbf{X} and parameters $\Theta = \{\Theta_1, \Theta_2\}$
 - ▶ Likelihood: $p(\mathbf{X}|\Theta)$
 - ▶ Prior: $p(\Theta)$
- ② Compute the posterior
 - ▶ Posterior:

$$p(\Theta|\mathbf{X}) = \frac{p(\mathbf{X}|\Theta)p(\Theta)}{p(\mathbf{X})}$$

- ▶ Can report summaries, e.g. posterior expectations

$$\mathbb{E}[h(\Theta)|\mathbf{X}]$$

and compute marginal posteriors

$$p(\Theta_1|\mathbf{X}) \text{ and } p(\Theta_2|\mathbf{X})$$

Section 2

- 1 Background
- 2 Variational Bayesian Inference

- 3 Hierarchical Dynamic Factor Analysis
Classical Factor Analysis
Dynamic Factor Analysis

Why Variational Bayesian Inference?

Consider a model with data \mathbf{X} and latent variables $\mathbf{Z} = \{\mathbf{Z}_1, \mathbf{Z}_2\}$.
The goal is to compute the joint posterior of the latent variables given
the data

$$p(\mathbf{Z}|\mathbf{X}) = \frac{p(\mathbf{X}|\mathbf{Z})p(\mathbf{Z})}{p(\mathbf{X})} \quad (1)$$

- ▶ Likelihood: $p(\mathbf{X}|\mathbf{Z})$
- ▶ Prior: $p(\mathbf{Z})$
- ▶ Evidence: $p(\mathbf{X})$

$$p(\mathbf{X}) = \int p(\mathbf{X}, \mathbf{Z}) d\mathbf{Z} \quad (2)$$

Why Variational Bayesian Inference?

For complex models (1) and (2) typically have no closed-form. As a result, the joint and marginal posteriors have to be approximated. Markov chain Monte Carlo (MCMC) has been the gold standard to solve this problem.

- ① Construct an ergodic Markov chain on \mathbf{Z} whose stationary distribution is the joint posterior $p(\mathbf{Z}|\mathbf{X})$
- ② Sample from the chain to collect samples from the stationary distribution
- ③ Approximate the posterior with an empirical estimate constructed from a subset of the collected samples
- ④ Use the subset of collected samples to estimate expectations of interest

Why Variational Bayesian Inference?

A major disadvantage for MCMC is that high levels of accuracy requires a great deal of time

- ① MCMC is typically slow, but ultimately accurate
- ② VB is typically much faster

Variational Bayesian Inference

Rather than use sampling, VB uses optimization to find an approximation to the posterior.

- ① Posit a family of “nice” approximate densities \mathcal{Q}
- ② Find a member of that family that is “closest” to the exact posterior, i.e.

$$q^*(\mathbf{Z}) = \arg \min_{q(\mathbf{Z}) \in \mathcal{Q}} \text{KL}(q(\mathbf{Z}) \parallel p(\mathbf{Z}|\mathbf{X})) \quad (3)$$

where

$$\text{KL}(q(\mathbf{Z}) \parallel p(\mathbf{Z}|\mathbf{X})) = \mathbb{E}_q[\log q(\mathbf{Z})] - \mathbb{E}_q[\log p(\mathbf{Z}|\mathbf{X})] \quad (4)$$

Variational Bayesian Inference

It can be shown that

$$\text{KL}(q(\mathbf{Z}) \parallel p(\mathbf{Z}|\mathbf{X})) = \mathbb{E}_q[\log q(\mathbf{Z})] - \mathbb{E}_q[\log p(\mathbf{X}, \mathbf{Z})] + \log p(\mathbf{X})$$

This reveals that the objective in (3) depends on the evidence, thus it cannot be computed directly.

Variational Bayesian Inference

Instead, we optimize an alternative objective that is equivalent to (3) up to an added constant called the **evidence lower bound** (ELBO)

$$\begin{aligned}\text{ELBO}(q) &= \mathbb{E}_q[\log p(\mathbf{X}, \mathbf{Z})] - \mathbb{E}_q[\log q(\mathbf{Z})] \\ &= \mathbb{E}_q\left[\log \frac{p(\mathbf{X}, \mathbf{Z})}{q(\mathbf{Z})}\right]\end{aligned}\tag{5}$$

It can be shown that

$$\arg \min_{q(\mathbf{Z}) \in \mathcal{Q}} \text{KL}(q(\mathbf{Z}) \parallel p(\mathbf{Z} | \mathbf{X})) = \arg \max_{q(\mathbf{Z}) \in \mathcal{Q}} \text{ELBO}(q)$$

Variational Bayesian Inference

The VB framework is now

- ① Posit a family of “nice” approximate densities \mathcal{Q}
- ② Find a member of that family that is “closest” to the exact posterior, i.e.

$$q^*(\mathbf{Z}) = \arg \max_{q(\mathbf{Z}) \in \mathcal{Q}} \text{ELBO}(q) \quad (6)$$

Solving this optimization problem is still difficult in general

- ▶ Using the mean-field assumption can make it easier

Mean-Field Assumption

- ① Partition the latent variables into M groups, say $\mathbf{Z}_1, \dots, \mathbf{Z}_M$
- ② Assume that the distributions in \mathcal{Q} factorize across the groups,
i.e.

$$\mathcal{Q} = \left\{ q : q(\mathbf{Z}) = \prod_{m=1}^M q_m(\mathbf{Z}_m) \right\}$$

- ③ Learning the optimal q now reduces to learning the optimal q_1, \dots, q_M
- ④ Straightforward to optimize via coordinate descent
- ⑤ This is **NOT** a modeling assumption

Mean-Field Assumption

Interestingly, under the mean-field assumption, the optimization problem for a single q_m has the solution:

$$q_m(\mathbf{Z}_m) = \frac{\exp\{\mathbb{E}_{q_{-m}}[\log p(\mathbf{X}, \mathbf{Z})]\}}{\int \exp\{\mathbb{E}_{q_{-m}}[\log p(\mathbf{X}, \mathbf{Z})]\} d\mathbf{Z}_m} \quad (7)$$

This establishes what is called the **coordinate ascent variational inference** algorithm.

Coordinate Ascent Variational Inference

Algorithm 1: Coordinate ascent variational inference

Input: Model $p(\mathbf{X}, \mathbf{Z})$, Data \mathbf{X} ,

Output: Variational density $q(\mathbf{Z}) = \prod_{m=1}^M q_m(\mathbf{Z}_m)$

Initialize: Variational densities $q_m(\mathbf{Z}_m)$

while the ELBO has not converged **do**

for $m \in \{1, 2, \dots, M\}$ **do**

 | Set $q_m(\mathbf{Z}) \propto \exp\{\mathbb{E}_{q_{-m}}[\log p(\mathbf{X}, \mathbf{Z})]\}$

 | **end**

 | Compute ELBO(q)

end

return $q(\mathbf{Z})$

Section 3

- 1 Background
- 2 Variational Bayesian Inference
- 3 Hierarchical Dynamic Factor Analysis
 - Classical Factor Analysis
 - Dynamic Factor Analysis

Classical Factor Analysis Model

The Classical (orthogonal) FA model assumes assumes the form

$$\mathbf{X} = \boldsymbol{\Lambda}\mathbf{F} + \boldsymbol{\epsilon}$$

where

- ① $\mathbf{X} = (X_1, \dots, X_N)^\top$ denotes the vector of observations
- ② $\boldsymbol{\Lambda} = [\lambda_{nk}]_{N \times K}$ denotes the matrix of factor loadings
- ③ $\mathbf{F} = (F_1, \dots, F_K)^\top$ denotes the vector of latent factors
- ④ $\boldsymbol{\epsilon} = (e_1, \dots, e_N)^\top$ denotes the vector of latent error terms

Normal Theory Assumptions

The Normal Theory Classical FA model assumes the form

$$\mathbf{X} = \boldsymbol{\Lambda}\mathbf{F} + \mathbf{e}$$

and adds the assumptions that

- ① $\mathbf{F} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_K)$
- ② $\mathbf{e} \sim \mathcal{N}(\mathbf{0}, \boldsymbol{\Sigma})$ where $\boldsymbol{\Sigma} = \text{diag}(\sigma_1^2, \dots, \sigma_N^2)$
- ③ F_k and e_n are independent for every pair k, n

Hierarchical Dynamic Factor Model

$$\mathbf{X}_{bst} = \boldsymbol{\Lambda}_{H.bs}(L)\mathbf{H}_{bst} + \mathbf{e}_{Xbst} \quad (8)$$

$$\mathbf{H}_{bst} = \boldsymbol{\Lambda}_{G.bs}(L)\mathbf{G}_{bt} + \mathbf{e}_{Hbst} \quad (9)$$

$$\mathbf{G}_{bt} = \boldsymbol{\Lambda}_{F.b}(L)\mathbf{F}_t + \mathbf{e}_{Gbt} \quad (10)$$

$$\boldsymbol{\Psi}_F(L)\mathbf{F}_t = \boldsymbol{\epsilon}_{Ft}, \quad (11)$$